

cse30 discussion 9 - final review

Ibrahim Awwal

July 27, 2015

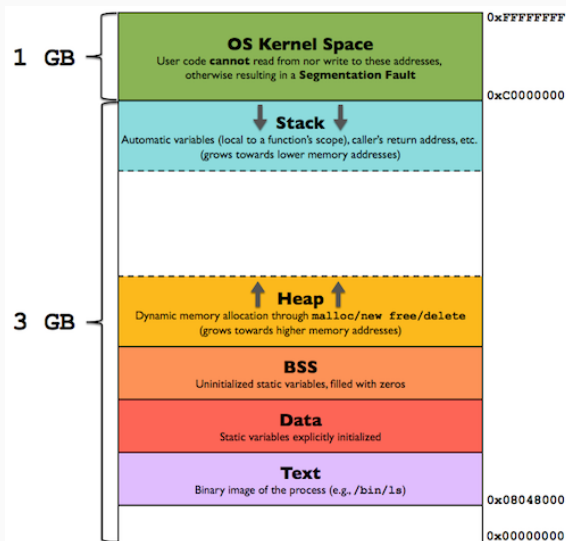
pointers and memory

pointer arithmetic

```
int arr[4] = {256, -1, 3, 0xAABBCCDD}  
int *ptr = arr;  
*(ptr+1);  
*(((char *) ptr) + 1);
```

- `ptr+1` increments by `1*sizeof(datatype)`
- Casting a pointer changes the increment
- Subtracting pointers within same array makes sense, adding does not - why?

heap/stack



- Processors read multiple bytes from memory at once
- Variable address must be aligned on a multiple of the size
- Struct fields should be ordered to minimize padding
- Array of Structs vs Struct of Arrays

structs in assembly

- Compute offsets into struct to read individual fields
- Be aware of alignment/padding when calculating offsets

```
struct foo{  
    char c;  
    short s;  
    int x;  
}
```

@ r0 is address of struct

```
ldrb r1, [r0, #0]    @ c
```

```
ldrh r2, [r0, #2]    @ s
```

```
ldr r3, [r0, #4]     @ x
```

multi dimensional vs multi level arrays

- Multi dimensional: `int X[3][4];`
- Multi level:
 - `int *X[3]; X[0] = (int *) malloc(4*sizeof(int));`
- Instructions needed to access an element?
- When would we use each type?

- Little Endian: LSB stored first
- Big Endian: MSB stored first
- x86, ARM are both Little Endian (though ARM can be configured Big Endian)
- Network byte order is Big Endian

Example: `int x = 0x12345678; &x=0x8000`

Address	0x8000	0x8001	0x8002	0x8003
Little	0x78	0x56	0x34	0x12
Big	0x12	0x34	0x56	0x78

assembly programming

- Clearing bits: BIC, Shifts, AND
- Setting bits: OR
- Masking

```
CMP r0, #0  
BEQ null  
LDR r0, [r0]  
null: MOV r0, #-1
```

1. Save temporary registers onto the stack (push)
 2. Put arguments in r0-r3 (additional arguments on stack)
 3. Branch and link to target address (bl <funcname>)
 4. Callee puts return value in r0-r1
 5. Callee branches back to link register (bx lr)
- Caller vs Callee save registers?
 - ARM Architecture Procedure Call Standard (AAPCS)

nested and recursive function calls

- Set up a function call as usual
- `bl <funcname>`
- After function returns, which instruction is executed next?

```
int some_func(int arg){  
    ...  
    return some_func(arg-1);  
}
```

- What does this recursive call correspond to?
- What about `return some_func(arg-1) + some_func(arg-2);`

- Basic operations: `push`, `pop`
- `sp` points to top of stack (last entry pushed)
- Stack grows downward
- Each function call has its own “stack frame”

- Not saving required registers onto the stack
- Using volatile registers between function calls
- Incorrect recursive calls (bl without returning)
- Conditional branches not accounting for fall through
- Not indexing by data type size
- Forgetting to increment loop counter

pipelining

- Idea: Split execution of an instruction into multiple stages
- Multiple instructions can then execute in parallel, also can shorten time for each stage
- Increases Instructions Per Cycle (IPC) or Instruction Level Parallelism (ILP)
- Less logic per stage \Rightarrow shorter clock period needed
- What is the drawback?
- What is the limitation of this technique?

5 stage pipeline

1. Instruction fetch
2. Instruction Decode/Register read
3. Execution
4. Data Memory
5. Register Write

- Not every set of instructions can be parallelized
- If one instruction uses result of previous instruction as operand, we need to wait for this result (data hazard)
- Register write happens in first half of cycle (stage 5)
- Register read happens in second half of cycle (stage 2)
- If we have a conditional branch, we can't continue pipelining instructions until we know whether the branch is taken (control hazard)

examples of data hazards

```
ADD r2, r0, r1
```

```
MUL r2, r2, r3
```

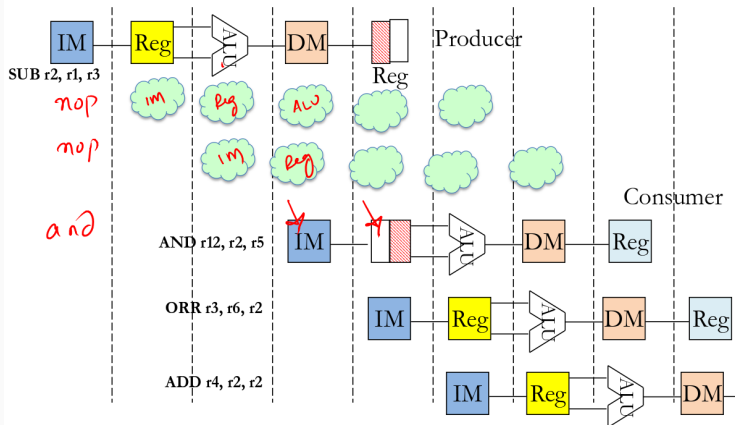
```
LDR r4, [r0]
```

```
AND r5, r4, r1
```

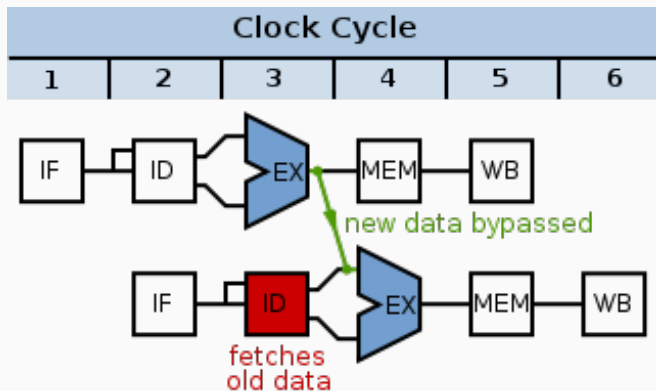
- Insert NOPs during compilation
- Stall pipeline (in hardware) and wait for result
- Reorder instructions (before assembling) to eliminate stalls
- Forward result from previous stages

Solving data hazards: s/w NOPs

2 instruction gap between producer and consumer



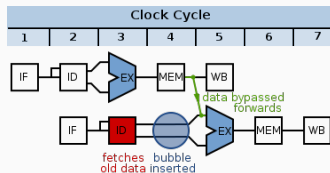
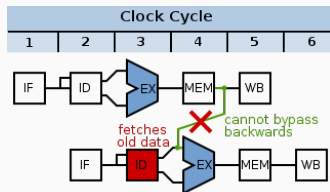
example: forwarding



examples: data memory

```
LDR r4, [r0]
```

```
AND r5, r4, r1
```



- Caches: Main memory can take 100s of cycles to access, so put commonly accessed data closer to CPU
- Branch Prediction: Predict whether a branch will be taken, where it will go
- Out of Order Execution: Reorder instructions to eliminate pipeline stalls
- Superscalar: Use multiple functional units to issue and process multiple instructions at once

practice final questions
